Eunsol CHOI
Assistant Professor at The University of Texas at Austin

○ https://eunsol.github.io/
↗ google scholar ↗ semantic scholar @ eunsol@utexas.edu
📍 2317 Speedway, University of Texas at Austin Austin, TX 78712 USA

EMPLOYMENT

August 2020-Now	Assistant Professor in the Department of Computer Science, The University of Texas at Austin
Sep 2019 - Aug 2020, 2022	Visiting Faculty Researcher, Google AI NY
Spring 2016	Research Intern, Google Mountain View
Summer 2014	Research Intern, University of Southern California
Summer 2012	Engineering Intern, Bloomberg
Fall 2010	Engineering Intern, Sentrana

EDUCATION

2019	Ph. D in Computer Science and Engineering, University of Washington
	Learning to Understand Entities in Text, Thesis Advisor: Luke Zettlemoyer, Yejin Choi
2012	Bachelor of Arts in Computer Science and Mathematics (both with distinction), Cornell University

PREPRINTS

Crafting In-context Examples according to LMs' Parametric Knowledge ARXIV, 2023
Yoonsang Lee*, Pranav Atreya*, Xi Ye, **Eunsol Choi** [pdf]

Clarify When Necessary: Resolving Ambiguity Through Interaction with LMs ARXIV, 2023
Michael J.Q Zhang, **Eunsol Choi** [pdf]

RECOMP: Improving Retrieval-Augmented LMs with Compression and Selective Augmentation ARXIV, 2023
Fangyuan Xu, Weijia Shi, **Eunsol Choi** [pdf]

Understanding Retrieval Augmentation for Long form Question Answering ARXIV, 2023
Hung-Ting Chen, Fangyuan Xu*, Shane Arora*, **Eunsol Choi** [pdf]

Complex Claim Verification with Evidence Retrieved in the Wild ARXIV, 2023
Jifan Chen, Grace Kim, Aniruddh Sriram, Greg Durrett, **Eunsol Choi** [pdf]

Unit-based Speech-to-Speech Translation Without Parallel Data ARXIV, 2023
Anuj Diwan, Anirudh Srinivasan, David Harwath, **Eunsol Choi** [pdf]

CONFERENCE / JOURNAL PUBLICATIONS

Mitigating Temporal Misalignment by Discarding Outdated Facts EMNLP, 2023
Michael J.Q. Zhang, **Eunsol Choi** [pdf]

Continually Improving Extractive QA via Human Feedback EMNLP, 2023
Ge Gao*, Hung-ting Chen*, Yoav Artzi, **Eunsol Choi** [pdf]

Propagating Knowledge Updates to LMs Through Distillation NEURIPS, 2023
hankar Padmanabhan, Yasumasa Onoe, Michael J.Q. Zhang, Greg Durrett, **Eunsol Choi** [pdf]

A Critical Evaluation of Evaluations for Long-form Question Answering ACL, 2023
Fangyuan Xu*, Yixiao Song*, Mohit Iyyer, **Eunsol Choi** [pdf]

Can LMs Learn New Entities from Descriptions? Challenges in Propagating Injected Knowledge ACL, 2023
Yasumasa Onoe, Michael J.Q Zhang, Shankar Padmanabhan, Greg Durrett, **Eunsol Choi** [pdf]

Concise Answers to Complex Questions: Summarization of Long-Form Answers ACL, 2023
Abhilash C. Potluri*, Fangyuan Xu*, **Eunsol Choi** [pdf]

When to Use Efficient Self Attention? Profiling Text, Speech and Image Transformer Variants ACL (SHORT), 2023
Anuj Diwan, **Eunsol Choi**, David Harwath [pdf]

Quantifying Train-Evaluation Overlap with Nearest Neighbors   Findings of ACL, 2023
Gauri Kambhatla, Thuy H. Nguyen, **Eunsol Choi** [pdf]

Continual Learning for On-Device Speech Recognition using Disentangled Conformers   ICASSP, 2023
Anuj Diwan, Ching-Feng Yeh, Wei-Ning Hsu, Paden Tomasello, **Eunsol Choi**, David Harwath, Abdelrahman Mohamed [pdf]

Understanding Postpartum Parents' Experiences via Two Digital Platforms   CSCW, 2023
Xuewen Yao, Miriam Mikhelson, Megan Micheletti, **Eunsol Choi**, S Craig Watkins, Edison Thomaz, Kaya De Barbaro [pdf]

Rich Knowledge Sources Bring Complex Knowledge Conflicts: Recalibrating Models to Reflect Conflicting Evidence   EMNLP, 2022
Hung-ting Chen, Michael J.Q. Zhang, **Eunsol Choi** [pdf]

Why is Winoground Hard? Investigating Failures in Visuolinguistic Compositionality   EMNLP, 2022
Anuj Diwan*, Layne Berry*, **Eunsol Choi**, David Harwath, Kyle Mahowald [pdf]

Generating Literal and Implied Subquestions to Fact-check Complex Claims   EMNLP, 2022
Jifan Chen, Aniruddh Sriram, **Eunsol Choi**, Greg Durrett [pdf]

Beyond Counting Datasets: A Survey of Multilingual Dataset Construction and Necessary Resources Findings of EMNLP, 2022
Xinyan Velocity Yu*, Akari Asai*, Trina Chatterjee, Junjie Hu, **Eunsol Choi** [pdf]

TyDiP: A Dataset for Politeness Classification in Nine Typologically Diverse Languages   Findings of EMNLP, 2022
Anirudh Srinivasan, **Eunsol Choi** [pdf]

Modeling Exemplification in Long-form Question Answering via Retrieval   NAACL, 2022
Shufan Wang, Fangyuan Xu, Laure Thompson, **Eunsol Choi**, Mohit Iyyer [pdf]

Entity Cloze By Date: Understanding what LMs know about unseen entities   Findings of NAACL (short), 2022
Yasumasa Onoe, Michael J.Q. Zhang, **Eunsol Choi**, Greg Durrett [pdf]

How Do We Answer Complex Questions: Discourse Structure of Long-form Answers   ACL, 2022
Fangyuan Xu, Junyi Jessi Li, **Eunsol Choi** [pdf]

Simulating Bandit Learning from User Feedback for Extractive Question Answering   ACL, 2022
Ge Gao, **Eunsol Choi**, Yoav Artzi [pdf]

Misinfo Reaction Frames: Reasoning about Readers' Reactions to News Headlines   ACL, 2022
Saadia Gabriel, Skyler Hallinan, Maarten Sap, Pemi Nguyen, Franziska Roesner, **Eunsol Choi**, Yejin Choi [pdf]

CREAK: A Dataset for Commonsense Reasoningover Entity Knowledge   NeurIPS Datasets and Benchmark, 2021
Yasumasa Onoe, Michael J.Q. Zhang, **Eunsol Choi**, Greg Durrett [pdf]

SituatedQA:Incorporating Extra-Linguistic Contexts into QA   [🏆OUTSTANDING PAPER] EMNLP, 2021
Michael J.Q. Zhang, **Eunsol Choi** [pdf]

Learning with Different Amount of Annotation: From Zero to Many Labels   EMNLP, 2021
Shujian Zhang, Chengyue Gong, **Eunsol Choi** [pdf]

Can NLI Models Verify QA Systems' Predictions?   Findings of EMNLP, 2021
Jifan Chen, **Eunsol Choi**, Greg Durrett [pdf]

Challenges in Information Seeking QA: Unanswerable Questions and Paragraph Retrieval   ACL, 2021
Akari Asai, **Eunsol Choi** [pdf]

Knowing About Questions Can Help: Improving Calibration in Question Answering   Findings of ACL, 2021
Shujian Zhang, Chengyue Gong, **Eunsol Choi** [pdf]

NeurIPS 2020 EfficientQA Competition: Systems, Analyses and Lessons Learned   JMLR, 2021
Sewon Min, Jordan Boyd-Graber, Chris Alberti, Danqi Chen, **Eunsol Choi**, Michael Collins, Kelvin Guu, Hannaneh Hajishirzi, Kenton Lee, Jennimaria Palomaki, Colin Raffel, Adam Roberts, Tom Kwiatkowski, and many workshop participants [pdf]

QED: A Framework and Dataset for Explanations in Question Answering TACL, 2021
Matthew Lamm, Jennimaria Palomaki, Chris Alberti, Daniel Andor, **Eunsol Choi**, Livio Baldini Soares, Michael Collins [pdf]

XOR QA: Cross-lingual Open-Retrieval Question Answering NAACL, 2021
Akari Asai, Jungo Kasai, Jonathan H Clark, Kenton Lee, **Eunsol Choi**, Hannaneh Hajishirzi [pdf]

Decontextualization: Making Sentences Stand-Alone TACL, 2021
Eunsol Choi, Jennimaria Palomaki, Matthew Lamm, Tom Kwiatkowski, Dipanjan Das, Michael Collins [pdf]

Entities as Experts: Sparse Memory Access with Entity Supervision EMNLP, 2020
Thibault Févry, Livio Baldini Soares, Nicholas FitzGerald, **Eunsol Choi**, Tom Kwiatkowski [pdf]

TYDI QA: A Benchmark for Information-Seeking Question Answering in Typologically Diverse Languages TACL, 2020
Jonathan H. Clark, **Eunsol Choi**, Michael Collins, Dan Garrette, Tom Kwiatkowski, Vitaly Nikolaev, Jennimaria Palomaki [pdf]

No Permanent Friends or Enemies: Tracking Dynamic Relationships Between Nations from News NAACL, 2019
Xiaochuang Han, **Eunsol Choi**, Chenhao Tan [pdf]

FlowQA: Grasping Time Flow in Conversational Machine Comprehension ICLR, 2019
Hsin-Yuan Huang, **Eunsol Choi**, Wen-tau Yih [pdf]

pair2vec: Compositional Word-Pair Embeddings for Cross-Sentence Inference NAACL, 2019
Mandar Joshi, **Eunsol Choi**, Omer Levy, Dan Weld and Luke Zettlemoyer [pdf]

Question Answering In Context EMNLP, 2018
Eunsol Choi*, He He*, Mohit Iyyer*, Mark Yatskar*, Percy Liang, Wen-tau Yih, Yejin Choi, Luke Zettlemoyer [pdf]

Neural Metaphor Detection In Context EMNLP (SHORT), 2018
Ge Gao, **Eunsol Choi**, Yejin Choi, Luke Zettlemoyer [pdf]

Ultra-Fine Entity Typing ACL, 2018
Eunsol Choi, Omer Levy, Yejin Choi, Luke Zettlemoyer [pdf]

Zero-Shot Relation Extraction via Reading Comprehension CoNLL, 2017
Omer Levy, Minjoon Seo, **Eunsol Choi** and Luke Zettlemoyer [pdf]

Truth of Varying Shades: Analyzing Language in Fake News and Political Fact-Checking EMNLP (SHORT), 2017
Hannah Rashkin, **Eunsol Choi**, Jin Yea Jang, Svitlana Volkova, Yejin Choi [pdf]

Coarse-to-Fine Question Answering for Long Documents ACL, 2017
Eunsol Choi, Daniel Hewlett, Jakob Uszkoreit, Illia Polosukhin, Alexandre Lacoste and Jonathan Berant [pdf]

TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension ACL, 2017
Mandar Joshi, **Eunsol Choi**, Dan Weld and Luke Zettlemoyer [pdf]

Document-level Sentiment Inference with Social, Faction, and Discourse Context ACL, 2016
Eunsol Choi, Hannah Rashkin, Luke Zettlemoyer and Yejin Choi [pdf]

Extracting Structured Scholarly Information from the Machine Translation Literature LREC, 2016
Eunsol Choi*, Matic Horvat*, Jonathan May, Kevin Knight and Daniel Marcu [pdf]

Scalable Semantic Parsing with Partial Ontologies ACL, 2015
Eunsol Choi, Tom Kwiatkowski, Luke Zettlemoyer [pdf]

Scaling Semantic Parsers with On-the-fly Ontology Matching EMNLP, 2013
Tom Kwiatkowski, **Eunsol Choi**, Yoav Artzi, Luke Zettlemoyer [pdf]

WORKSHOP PROCEEDINGS

MIA 2022 Shared Task: Evaluating Cross-lingual Open-Retrieval Question Answering for 16 Diverse Languages NAACL 2022
WORKSHOP
Akari Asai, S. Longpre, Jungo Kasai, Chia-Hsuan Lee, Rui Zhang, Junjie Hu, Ikuya Yamada, J. Clark, **Eunsol Choi** [pdf]

MRQA 2019 Shared Task: Evaluating Generalization in Reading Comprehension EMNLP 2019 Workshop
Adam Fisch, Alon Talmor, Robin Jia, Minjoon Seo, **Eunsol Choi**, Danqi Chen [pdf]

Hedge detection as a lens on framing in the GMO debates: A position paper ACL 2012 Workshop on Extra-Propositional Aspects of Meaning in Computational Linguistics
Eunsol Choi, Chenhao Tan, Lillian Lee, Danescu-Niculescu-Mizil, Jennifer Spindel [pdf]

Academic Service

Faculty Recruiting	UT CS-2021, UT Moody College of Journalism 2021
Graduate Student Recruiting	UT CS-2020, UT CS-2023
Turing Scholars Recruiting	UT CS-2022
NSF Panel	December 2021, June 2023
Program Chair	AKBC 2022
Workshop Chair	ACL 2024
Publicity Chair	EMNLP 2022
Workshop Chair	AKBC 2021
Senior Area Chair	(Question Answering) EMNLP 2021, ACL 2023, (Multilinguality) EMNLP 2022
Junior Area Chair	(Question Answering) EMNLP 2019 NAACL 2021, (Semantics - Sentence Level) ACL 2020, (Information Extraction) EMNLP 2020, ICLR 2021-2, (Datasets and Benchmarks) NeurIPS 2021-2, ICML 2023, ICLR 2024
Action Editor	ARR
Standing Reviewer	TACL, TMLR
Reviewer	ACL 2015-8, NAACL 2018, NeurIPS 2018-9, ICML 2020, ICLR 2021, EMNLP 2015-8, *SEM 2018, EACL 2017, NAACL AKBC Workshop 2016, NeurIPS AKBC Workshop 2017, ACL RoboNLP Workshop 2017, ACL Rep4NLP Workshop 2019
Workshop	Student Research Workshop, NAACL 2016, ACL19
Co-Organizer	Workshop on Machine Reading for Question Answering, ACL18, EMNLP19, EMNLP21 [website]
	EfficientQA Workshop, NeurIPS 2020 [website]
	Workshop on Multilingual Information Access (MIA), NAACL 2022 [website]
	Workshop on Multilingual Information Access (MIA), NAACL 2022 [website]
	Optimism, Pessimism, and Realism of Large Language Models, Korea [website]
	Emerging Generalization Settings, Simons Institute Workshop 2024 [website]

Student Advising

Ph.D Supervisor	Jifan Chen (co-supervised with Greg Durrett), 2023 → Amazon AWS
	Yasumasa Onoe (co-supervised with Greg Durrett), 2023, → Google AI
Master's Thesis Advisor	Rishi Salem (co-supervised with David Harwath), 2023
	Anirudh Srinivasan, 2023 → Tenstorrent Inc.
	Hung-Ting Chen, 2023 → UT Austin Ph.D
Undergraduate Advisor	Abhilash C. Potluri, supervisor, 2022 → UIUC Masters
	Pranav Atreya, supervisor, 2023 → UC Berkeley Ph.D

Teaching

Fall 2022	UT Austin CS 388: Natural Language Processing (grad) [website]
Spring 2021-3	UT Austin CS 378: Natural Language Processing (undergrad) [website]
Fall 2020,3	UT Austin CS 395T: Topics in Natural Language Processing [website]
June 2023	High School Summer Camp at UT CS, Academy for All, Academy for Women [website]

INVITED TALK

Semantic parsing with Freebase: a very large, yet wildly incomplete knowledge base

NAACL 2016, Human-Computer QA Workshop

Scaling up Reading Comprehension

Nov 2017, Stanford NLP Seminar, Salesforce Research [website]

Contextual Machine Reading

Sept 2018, Ohio State University, Heidelberg University [website]

Scalable and Interactive Machine Reading

Nov 2018, Seoul National University [website]

Learning to Understand Entities in Text

Jan-Apr 2019, KAIST, UIUC, FAIR, Google, UBC, U of Waterloo, Georgia Tech, UT Austin, UNC, UC Irvine, USC, Yale, EPFL, UMass Amherst, CU Boulder

Learning to Understand Language In Context

Aug. 20, Seoul National University, Google India AI Summer Camp, EMNLP Workshop on Computational Approaches to Discourse

Addressing Information Seeking Queries

Feb 2021, UC Santa Barabara, Johns Hopkins NLP Seminar [slides]

Context for Interpreting Benchmark Performances

Aug 2021, ACL Workshop on Benchmarking: Past, Present and Future [youtube]

Addressing Information Seeking Queries: From Finding to Presenting Answers

Aug 2021, Seoul National University

Incorporating extra-linguistic contexts and entity knowledge into natural language processing

Oct 2021, The 2nd Workshop on Unstructured and Structured KBs [website]

How Machines Understand Text and Where They Fail

Dec 2021, Ewha Womans University

Information Seeking Language Understanding in a Dynamic World

May 2022, Stanford NLP seminar, Rice University, ACL Spotlight Talks for Young Rising Stars [link] [youtube]

Knowledge-rich NLP models in a dynamic real world

Nov 2022, EvoNLP workshop EMNLP 2022, UW, AI2, UT MLL Research Symposium

Adapting Language Models to Dynamic World

May 2023, University of Maryland, Johns Hopkins University, UMass Amherst [slides]

Keynote at Korean Artificial Intelligence Graduate School Council Symposium

Aug 2023, Seoul

Knowledge Augmentation for Language Models

Sept 2023, Bar Ilan University, Tel Aviv University, KUIS AI talk series [youtube]

Knowledge Augmentation for Language Models

Nov 2023, Cornell Tech, NYU [slides]

AWARDS / HONORS

Spotlight Talks for Young Rising Stars - ACL	2022
Outstanding Paper Award - EMNLP	2021
AAAI New Faculty Highlights	2021
Participant, Heidelberg Laurete Forum	2019
Facebook (Meta) Graduate Research Fellowship	2018-2019
Rising Stars in EECS Workshop	2017
Honorable Mention, CRA Outstanding Undergraduate Research Award	2011

GRANTS

Research grant	NSF, RI: Medium: Multilingual Long-form QA with Retrieval-Augmented Language Models, 2023 (Co-PI: Mohit Iyyer), $1.2M
	Gift, Question Answering over Heterogenous Data Sources, Home Depot, 2022-3, $60K
	Open Philanthropy, 2022 (Co-PI: Greg Durrett, Raymond Mooney), $432K total
	UT Machine Learning Laboratory 2021/2 (solo), 2022/3 (Co-PI: David Harwath), $60K each
	Google Research Award, Answering information seeking questions in the wild, 2021, $60K
	UT Good Systems (Tier II) 2020 Co-PI: Kaya de Barbaro, Craig Watkins, James Pennebaker, $98K total
Sponsored Research	CISCO, Cascaded QA Model for Efficient Inference, 2023-4, $200K
	Sony AI, Externalizing and Internalizing Knowledge in Large Language Models, 2023 (Co-PI: Greg Durrett), $100K total