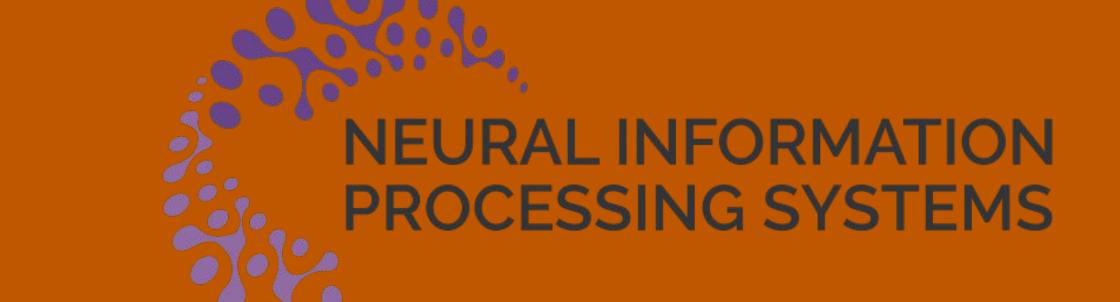


Propagating Knowledge Updates to LMs Through Distillation

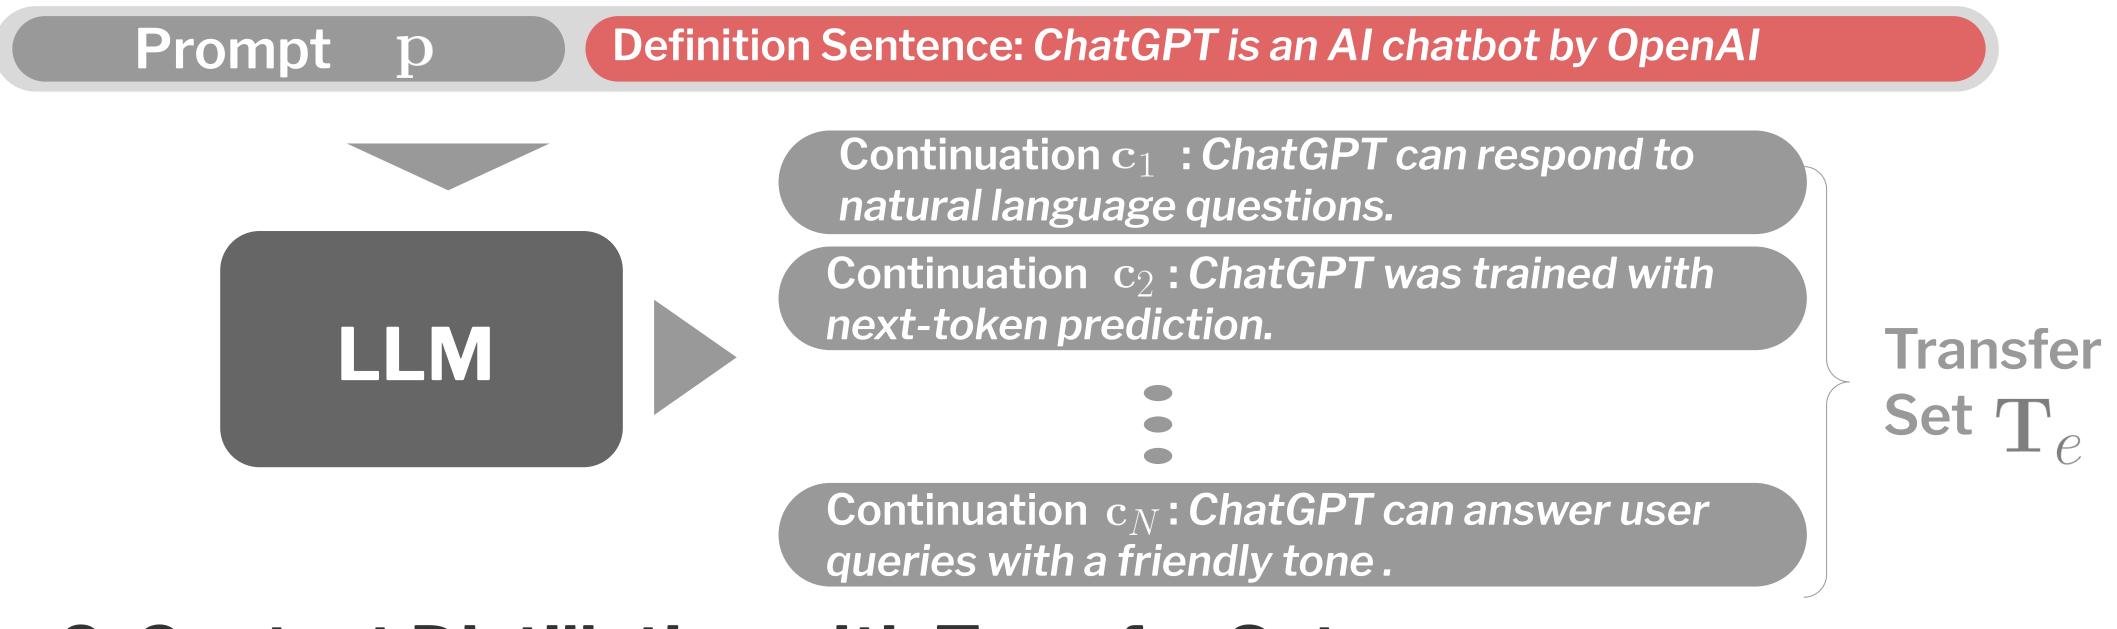


Shankar Padmanabhan, Yasumasa Onoe, Michael J.Q. Zhang, Greg Durrett, Eunsol Choi

Method

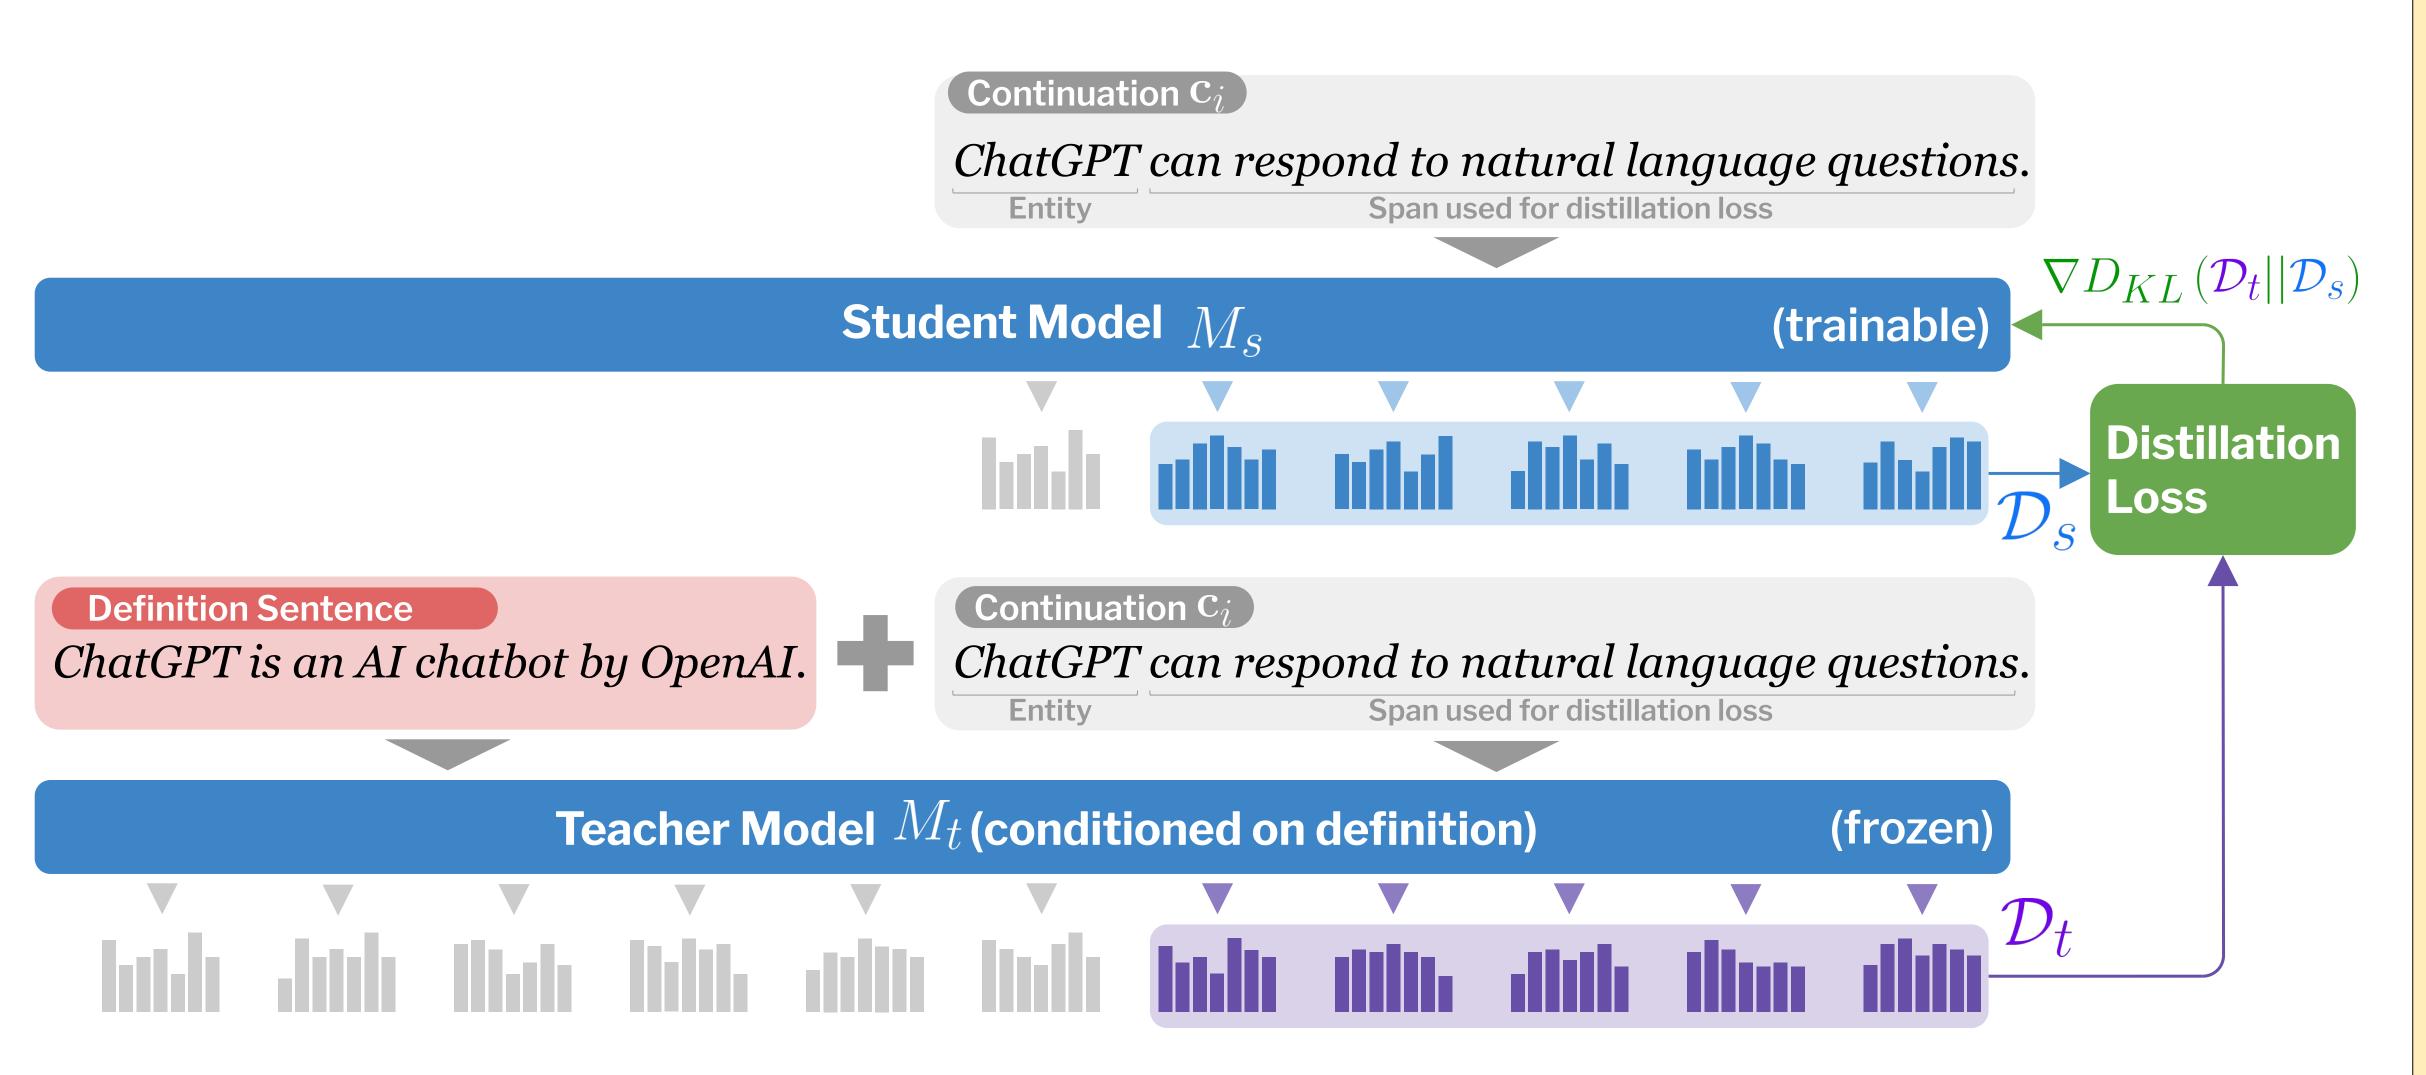
Step 1: Transfer Set Generation

Transfer set: A set of sentences that represent potential inferences given the definition statement. These are continuations of the definition sampled from the language model.



Step 2: Context Distillation with Transfer Set

The language model (student) is trained to match its distribution when the same model is conditioned on the definition (teacher).



Experimental Setup

Datasets:

Entity Inferences (Onoe et al., 2023)
Manually crafted probe sentences using templates
Measure accuracy on cloze span
Definition: Hurricane Nana was a minimal Category 1 hurricane that caused moderate damage across Belize in early September 2020.
Sentence: Hurricane Nana (2020) totally [MASK] my house.
Entity: Hurricane Nana (2020) totally [MASK] my house.
Entity: mRNA vaccine uses a copy of a molecule called messenger RNA to produce an immune response.
Sentence: mRNA vaccines do not affect or reprogram [MASK].
Entity: mRNA vaccine
Year: 2020
Label: DNA inside the cell

Evaluation Metrics

- Edit Efficacy: The success of the model edit at propagating the injected knowledge
- Edit Specificity: The updated models performance on queries about unrelated entities. Ideally, this should not change post-update.

Task: Knowledge Propagation

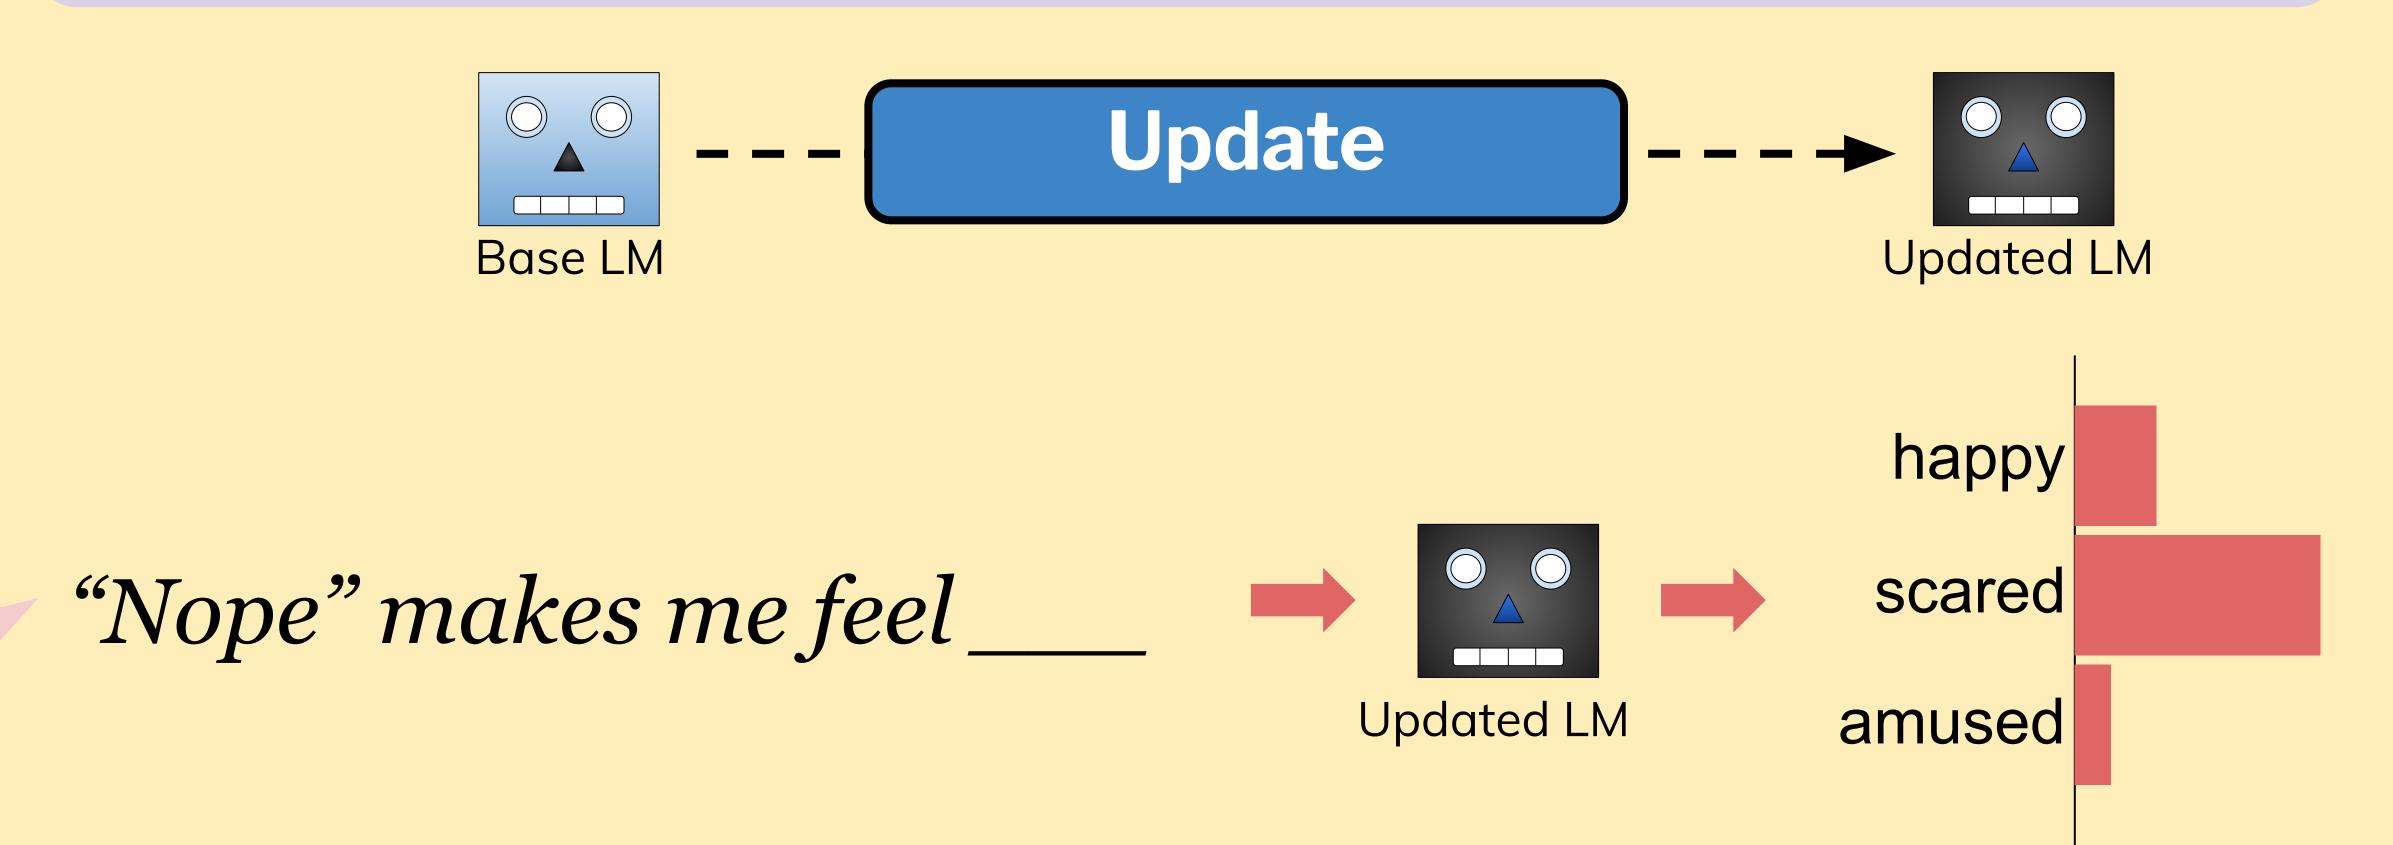
Prior work edits the parameters of LMs to update their knowledge ("Rishi Sunak is now the British PM") and evaluates these updates directly ("Who is the PM of the UK?"). We focus on **propagation** of injected knowledge, testing whether LMs can make inferences based on injected facts.

We develop a distillation-based knowledge editing method that can propagate injected knowledge!

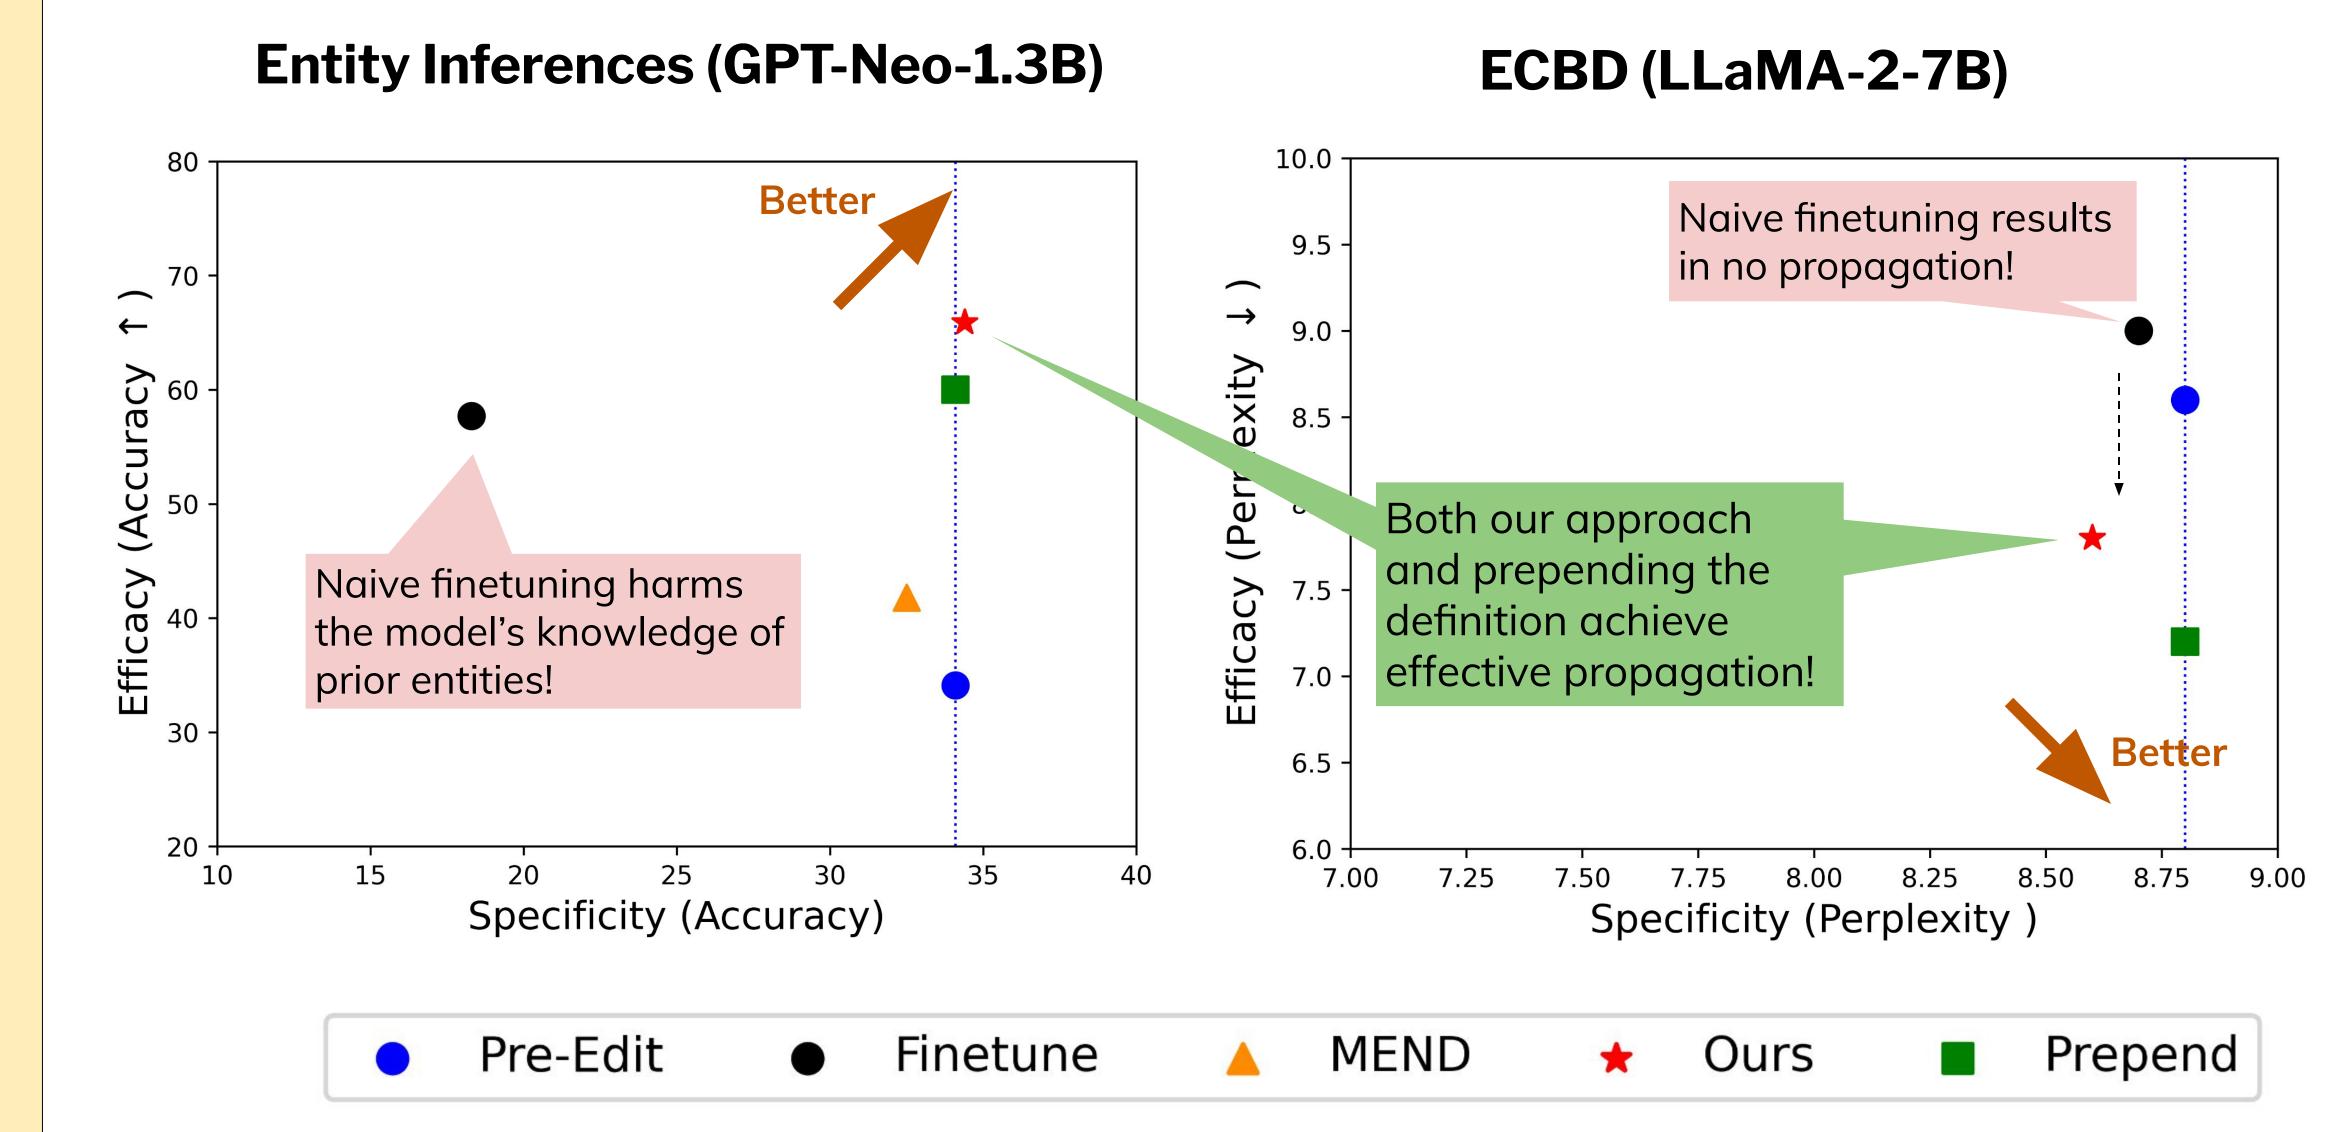
We **update** an LM on a definition sentence of a **new entity** using any knowledge editing method such as finetuning, MEND, or ROME.

The updated LM is evaluated for propagation of the knowledge: can the model make inferences based on the fact?

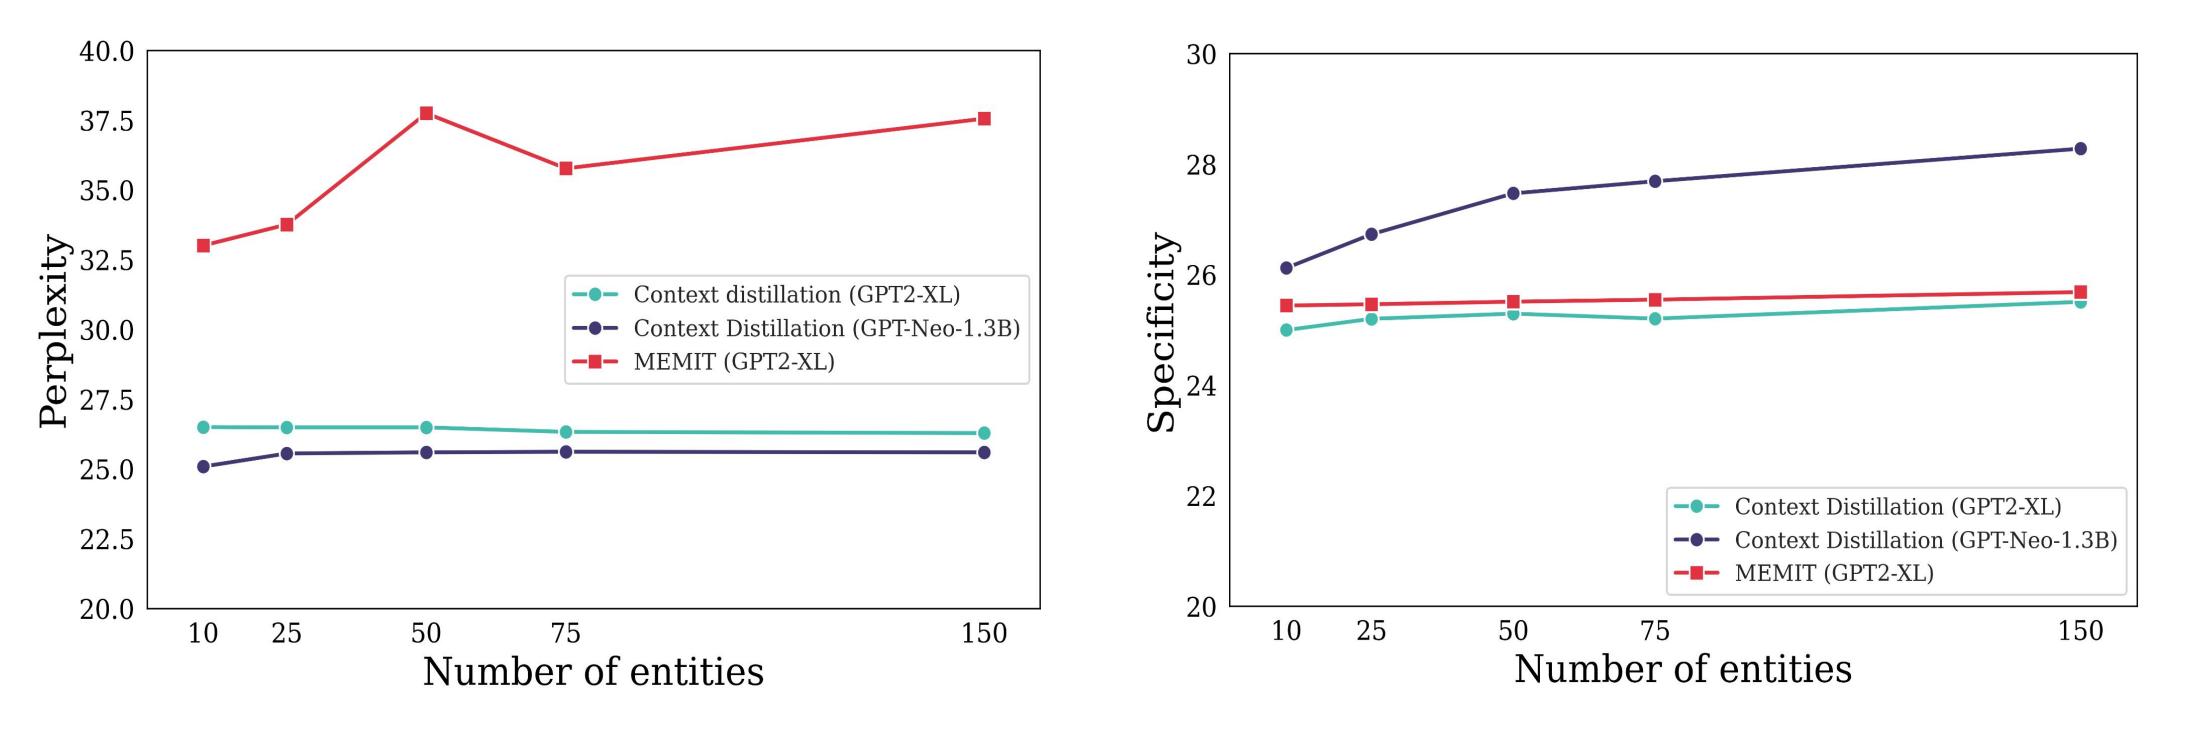
"Nope" is a 2022 American Western science fiction horror film written, directed, and produced by Jordan Peele.



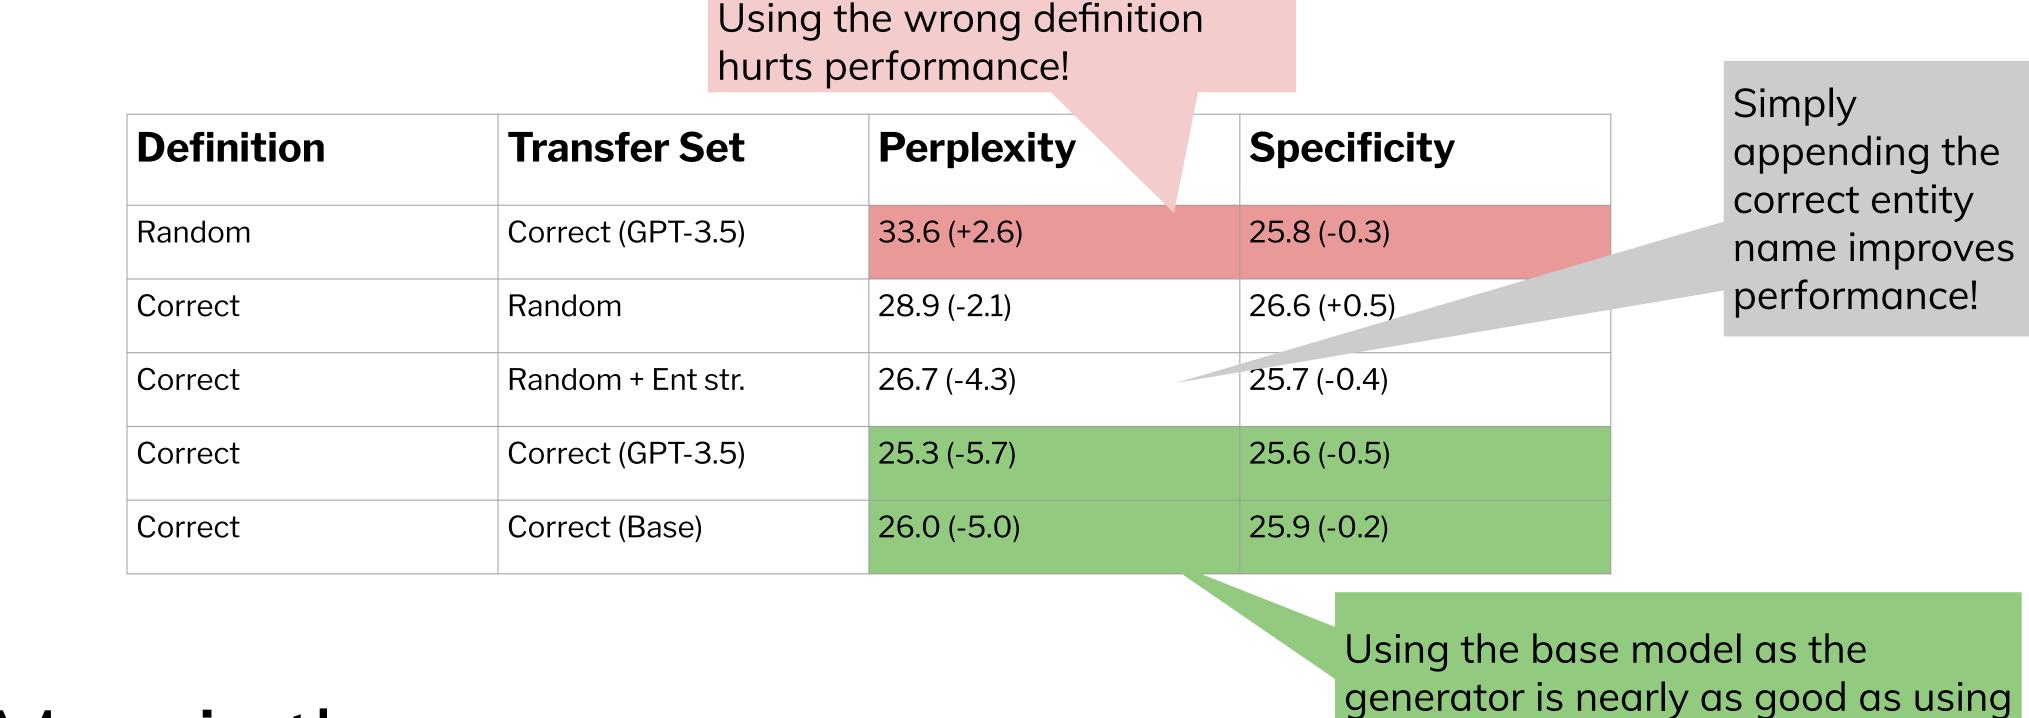
Results



Our method scales to many entities at once! (ECBD)



How much does the definition or transfer set matter?



- More in the paper:
 - Does the model memorize the definition during distillation?

SOTA model!

- What kind of updates is distillation capable of making?
- How does our method perform on Counterfact? Propagates to related info, but not as good as ROME. (However, ROME performs very poorly on ECBD.)

Code available at github.com/shankarp8/knowledge_distillation